

Performance-Evaluation der potenziellen Engpässe der neuronalen Netzimplementierung in SetlX

(network.stlx durch sigmoid_timing.stlx oder zip_timing.stlx ersetzen zum Testen)

1. Aufruf der Funktion zip(l1, l2). Bildung jeweils einer Liste für l1 und l2 (durch toList(v)) bedeutet zusätzlichen Rechenaufwand.

Codeabschnitte mit Zeitmessungspunkten:

```
sgd := procedure(training_data, epochs, mini_batch_size, eta, test_data) {
  if(test_data != null) {
    n_test := #test_data;
  }
  n := #training_data;

  for(j in {0..epochs}) {
    s1 := now();
    training_data := shuffle(training_data);
    // Get mini-batches from the training data to train the network
    mini_batches := [ training_data[k..k+mini_batch_size-1] : k in [1,mini_batch_size..n] ];

    for(mini_batch in mini_batches) {
      update_mini_batch(mini_batch, eta);
    }
    epoche_time := now() - s1;
    // Visual output
    if(test_data != null) {
      ev := evaluate(test_data);
      //print("Epoch $j$: $ev$ / $n_test$");
      print("Zipping-time:\t" + zip_time);
      print("Epoche-time:\t" + epoche_time);
      print("--> " + 100.0 * zip_time/epoche_time + "%");
      this.zip_time := 0;
    }
    else {
      print("Epoch $j$ complete");
    }
  }
};
```

```
/* Computes the Cartesian product of two matrices or vectors */
zip := procedure(l1, l2) {
  s1 := now();
  res := toList(l1) >< toList(l2);
  this.zip_time += (now() - s1);

  return res;
};
/* casts vector to list */
toList := procedure(v) {
  return [v[i] : i in [1..#v]];
};
```

Anzahl Datensätze: 10.000 Testsätze, 10.000 Trainingssätze

Rechnerdaten: Intel Core i7-4720HQ, 16GB RAM

Ergebnisse 1.:

```
Start SGD
Zipping-time: 7154
Epoche-time: 23136
--> 30.921507607192254%
Zipping-time: 6451
Epoche-time: 19665
--> 32.8044749555047%
Zipping-time: 6006
Epoche-time: 20937
--> 28.68605817452357%
Zipping-time: 6371
Epoche-time: 19349
--> 32.926766241149416%
Zipping-time: 6550
Epoche-time: 20508
--> 31.938755607567778%
Zipping-time: 6229
Epoche-time: 18742
--> 33.23551381922954%
Zipping-time: 6248
Epoche-time: 18913
--> 33.03547824247872%
```

2. sigmoid_prime(z) und sigmoid_vector(z)

Codeabschnitte mit Zeitmessungspunkten:

```
sgd := procedure(training_data, epochs, mini_batch_size, eta, test_data) {
  if(test_data != null) {
    n_test := #test_data;
  }
  n := #training_data;

  for(j in {0..epochs}) {
    s1 := now();
    training_data := shuffle(training_data);
    // Get mini-batches from the training data to train the network
    mini_batches := [ training_data[k..k+mini_batch_size-1] : k in [1,mini_batch_size..n] ];

    for(mini_batch in mini_batches) {
      update_mini_batch(mini_batch, eta);
    }
    epoche_time := now() - s1;
    // Visual output
    if(test_data != null) {
      ev := evaluate(test_data);
      //print("Epoch $j$: $ev$ / $n_test$");
      print("Sigmoid-time:\t" + sigmoid_time);
      print("Epoche-time:\t" + epoche_time);
      print("--> " + 100.0 * sigmoid_time/epoche_time + "%");
      this.sigmoid_time := 0;
    }
    else {
      print("Epoch $j$ complete");
    }
  }
};
```

```
// Sigmoid function for vectors
// 1.0/(1.0+np.exp(-z))
sigmoid_vector := procedure(z) {
  // z is a vector, so the function has to be used on every part of it
  s1 := now();
  res := la_vector([ 1.0/(1.0 + exp(- z[i] )) : i in [1..#z] ]);
  this.sigmoid_time += (now() - s1);
  return res;
};

// Derivative of the sigmoid function, when z is a vector
// sigmoid(z)*(1-sigmoid(z))
sigmoid_prime := procedure(z) {
  s := sigmoid_vector(z);
  s1 := now();
  res := la_matrix([ [ s[i] * (1 - s[i]) ] : i in [1..#s] ]);
  this.sigmoid_time += (now() - s1);
  return res;
};
```

Ergebnisse 1.:

```
Start SGD
Sigmoid-time: 1423
Epoche-time: 23065
--> 6.169520919141556%
Sigmoid-time: 1391
Epoche-time: 20955
--> 6.63803388212837%
Sigmoid-time: 1328
Epoche-time: 24388
--> 5.445300967689027%
Sigmoid-time: 1220
Epoche-time: 21434
--> 5.691891387515163%
Sigmoid-time: 1583
Epoche-time: 22733
--> 6.963445211806625%
Sigmoid-time: 1337
Epoche-time: 20903
--> 6.396211070181313%
```